


# Neural Networks Basics

Dai Bui

# Simple Neural Network



$$y = W^T x + b$$

# Simple Neural Network

- Consider the simple case: we have sample data  $\{x, y\}$  where  $x$  is input to a real world object and  $y$  is the measured output of the object.
- Let  $f(y; W)$  be some objective function that we want to minimize
- We want to train the neural network, e.g., find  $W$  matrix, so that the predicted output  $\hat{y}$  is close to measured  $y$ .

$$\|f(\hat{y}; W) - f(y; W)\|^2 \approx 0$$

- The idea is to adjust weight matrix  $W$  to reduce the error


# Neural Networks Training

- Neural networks training is composed of two phases
  - Forward propagation:  $\hat{y} = W^T x + b$
  - Backward propagation
    - Compute error:  $L(W) = \|f(\hat{y}; W) - f(y; W)\|^2$
    - Propagate the error  $e$  back to **adjust**  $W$  using the gradient descent method

$$W = W - \alpha \frac{\partial L}{\partial W}$$

# Backpropagation


- Let assume that we can compute the gradient:  $\frac{\partial f(y;W)}{\partial y}$
- Now we need to compute:
  - $\frac{\partial L}{\partial W}$  : gradient is used to adjust  $W$
  - $\frac{\partial L}{\partial x}$  : gradient is propagated back to adjust weight matrices of **previous layers** in case we have multi-layer network



# Backpropagation Intuition

- We have the following relation: if  $y = f(x)$  then  $x \rightarrow x + \Delta x \Rightarrow y \rightarrow \approx y + \frac{\partial y}{\partial x} \Delta x$
- Let us consider a change  $\Delta x_0$ , this change will lead to:

$$\begin{bmatrix} y_0 \\ y_1 \\ y_2 \end{bmatrix} \rightarrow \begin{bmatrix} y_0 + w_{00}\Delta x_0 \\ y_1 + w_{01}\Delta x_0 \\ y_2 + w_{02}\Delta x_0 \end{bmatrix} = y + \Delta x_0 \begin{bmatrix} w_{00} \\ w_{01} \\ w_{02} \end{bmatrix}$$
$$\Rightarrow f(y) \rightarrow f(y) + \left[ \frac{\partial L}{\partial y} \right]^T \Delta x_0 \begin{bmatrix} w_{00} \\ w_{01} \\ w_{02} \end{bmatrix}$$



# Backpropagation Intuition

- In short, a change  $\Delta x_0$  will lead to an approximate change of  $\left[\frac{\partial L}{\partial y}\right]^T \Delta x_0 \begin{bmatrix} w_{00} \\ w_{01} \\ w_{02} \end{bmatrix}$  in the objective function

$$\frac{\partial L}{\partial x_0} = \lim_{\Delta x_0 \rightarrow 0} \frac{\left( L(y) + \left[\frac{\partial L}{\partial y}\right]^T \Delta x_0 \begin{bmatrix} w_{00} \\ w_{01} \\ w_{02} \end{bmatrix} \right) - L(y)}{\Delta x_0} = \left[\frac{\partial L}{\partial y}\right]^T \begin{bmatrix} w_{00} \\ w_{01} \\ w_{02} \end{bmatrix} = \frac{\partial L}{\partial y_0} w_{00} + \frac{\partial L}{\partial y_1} w_{01} + \frac{\partial L}{\partial y_2} w_{02}$$

# Backpropagation Mathematically

- By the chain rule, we know

$$\bullet \frac{\partial L}{\partial x_0} = \frac{\partial L}{\partial y} \frac{\partial y}{\partial x_0} = \begin{bmatrix} \frac{\partial L}{\partial y_0} \\ \frac{\partial L}{\partial y_1} \\ \frac{\partial L}{\partial y_2} \end{bmatrix}^T \begin{bmatrix} w_{00} \\ w_{01} \\ w_{02} \end{bmatrix} = \frac{\partial L}{\partial y_0} w_{00} + \frac{\partial L}{\partial y_1} w_{01} + \frac{\partial L}{\partial y_2} w_{02}$$



# Backpropagation Input Gradient

$$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial y} W^T$$


# Backpropagation Weight Gradient

$$\frac{\partial L}{\partial W} = X^T \frac{\partial L}{\partial y}$$

# Gradient Descent


$$W = W - \alpha \frac{\partial L}{\partial W}$$

# Multilayer Neural Networks



# Regularization

- In case we use too many layers as well as too **wide** hidden layers, it is very easy to fit the data when training but prediction is often wrong. This is the **overfitting** problem
- To mitigate the problem, we add a regularization factor to reduce the number of weights used.



$$L(W) = f(x; W) + \lambda R(W)$$

$$R(W) = \sum_{i,j} w_{ij}^2$$

# Vectorizing Across Multiple Samples

- When we feed one sample at a time into a network of  $l$  hidden layers, the computation is composed of multiple vector-matrix multiplication

$$y = ((xW_1 + b_1) \dots W_l) + b_l$$


- In general, it is much faster if we can combine multiple vector-matrix multiplications into one single matrix-matrix multiplications

$$r_0 = v_0M, r_1 = v_1M, \dots, r_n = v_nM \Rightarrow \begin{bmatrix} r_0 \\ r_1 \\ \vdots \\ r_l \end{bmatrix} = \begin{bmatrix} v_0 \\ v_1 \\ \vdots \\ v_l \end{bmatrix} M$$

- We want to combine input multiple samples together similarly when training to speed up computation
- We call this group of samples **batch**.

# Softmax Classifier

- Suppose that we have the following output after the layers



- How do we interpret those raw score numbers to determine the prediction outcome?

$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

- Softmax value is the normalized probability assigned to the correct label  $y_i$  given input  $x_i$

# Cross Entropy

- Information theory: The cross-entropy between a **true** distribution  $p$  and an **estimated** distribution  $q$  is defined as:

$$H(p, q) = - \sum_x p(x) \log q(x)$$

- Because the true distribution  $p = [0, \dots, 1, \dots, 0]$  where 1 is at the  $y_i$ -th position and  $q(y_i) = P(Y = y_i | X = x_i) = \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$ , therefore the cross entropy between the “predicted” and the “true” distributions becomes

$$\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}} = -s_{y_i} + \log \sum_j e^{s_j}$$



# Cross Entropy Loss

- From information theory:

$$H(p, q) = H(p) + D_{KL}(p||q)$$

- $H(p)$  is entropy of true distribution  $p$ , which is 0
- $D_{KL}(p||q)$  is Kullback–Leibler divergence, which is a measure how  $p$  is different from  $q$
- As we want  $q$  to become similar to  $p$ , we need to **minimize** the cross entropy, in other words, we want to minimize the loss function

$$L_i = -s_{y_i} + \log \sum_j e^{s_j}$$

# Cross Entropy Loss

- Notice that

$$L_i = -\log P(Y = y_i | X = x_i)$$

- Minimizing cross entropy loss function is the same as **maximizing** the probability of the correct class  $P(Y = y_i | X = x_i)$

# Gradient of Cross Entropy Loss

- Given input samples  $m$  input samples  $\{x^j, y^j\}$
- For each sample, we have the loss function, where  $C$  is the number classes, or possible outcomes

$$\begin{aligned} L(W) &= - \sum_{i=1}^C y_i^j \log \left( \frac{e^{s_i}}{\sum_{k=1}^C e^{s_k}} \right) \\ &= - \sum_{i=1}^C y_i^j \log \left( \frac{e^{W_i^T x^j}}{\sum_{k=1}^C e^{W_k^T x^j}} \right) \\ &= - \sum_{i=1}^C y_i^j \log \left( e^{W_i^T x^j} \right) - y_i^j \log \left( \sum_{k=1}^C e^{W_k^T x^j} \right) \\ &= - \sum_{i=1}^C y_i^j W_i^T x^j + \log \left( \sum_{k=1}^C e^{W_k^T x^j} \right) \end{aligned} \quad \text{##notice that } \sum_i^C y_i^j = 1$$

# Gradient of Cross Entropy Loss

$$\begin{aligned}\frac{\partial L(W)}{\partial W_i} &= -y_i^j x^j + \frac{e^{W_i^T x^j}}{\sum_{k=1}^C e^{W_k^T x^j}} \\ &= -y_i^j x^j + p_i^j x^j \\ &= (p_i^j - y_i^j) x^j\end{aligned}$$

As a result

$$\frac{\partial L(W)}{\partial W} = \begin{bmatrix} (p_0^j - y_0^j)(x^j)^T \\ (p_1^j - y_1^j)(x^j)^T \\ \vdots \\ (p_n^j - y_n^j)(x^j)^T \end{bmatrix} = \begin{bmatrix} (p_0^j - y_0^j) \\ (p_1^j - y_1^j) \\ \vdots \\ (p_n^j - y_n^j) \end{bmatrix} (x^j)^T$$

# Batch Gradient of Cross Entropy Loss

- Because we often compute loss function for batch, given

$$X = \begin{bmatrix} (x^1)^T \\ (x^2)^T \\ \vdots \\ (x^m)^T \end{bmatrix}, Y = \begin{bmatrix} (y^1)^T \\ (y^2)^T \\ \vdots \\ (y^m)^T \end{bmatrix}, P = \begin{bmatrix} (p^1)^T \\ (p^2)^T \\ \vdots \\ (p^m)^T \end{bmatrix}$$

then we have

$$\frac{\partial L(W)}{\partial W} = \frac{1}{m} X^T (P - Y)$$

# What about the Gradient of Regularization?

$$\frac{\partial L(W)}{\partial W} = \frac{1}{m} X^T (P - Y) + 2\lambda W$$

# Stable Softmax

- What happens if  $s_j$  in  $\frac{e^{s_k}}{\sum_j e^{s_j}}$  are big?
  - The exponential function can be numerically overflowed
- Notice that:  $\frac{e^{s_k}}{\sum_j e^{s_j}} = \frac{e^{-c} e^{s_k}}{e^{-c} \sum_j e^{s_j}} = \frac{e^{s_k - c}}{\sum_j e^{s_j - c}}$
- We can basically subtract a constant from each score without changing the softmax probability
- To avoid the case that some  $s^j$  is too big, we do the following

$$s_j = s_j - \max_k s_k$$