


Introduction to Convolutional Neural Networks

Dai Bui


Intuition of Convolution

- Convolution, similar to cross correlation of two signals is high when they look similar



Intuition for CNN for Pattern Recognition

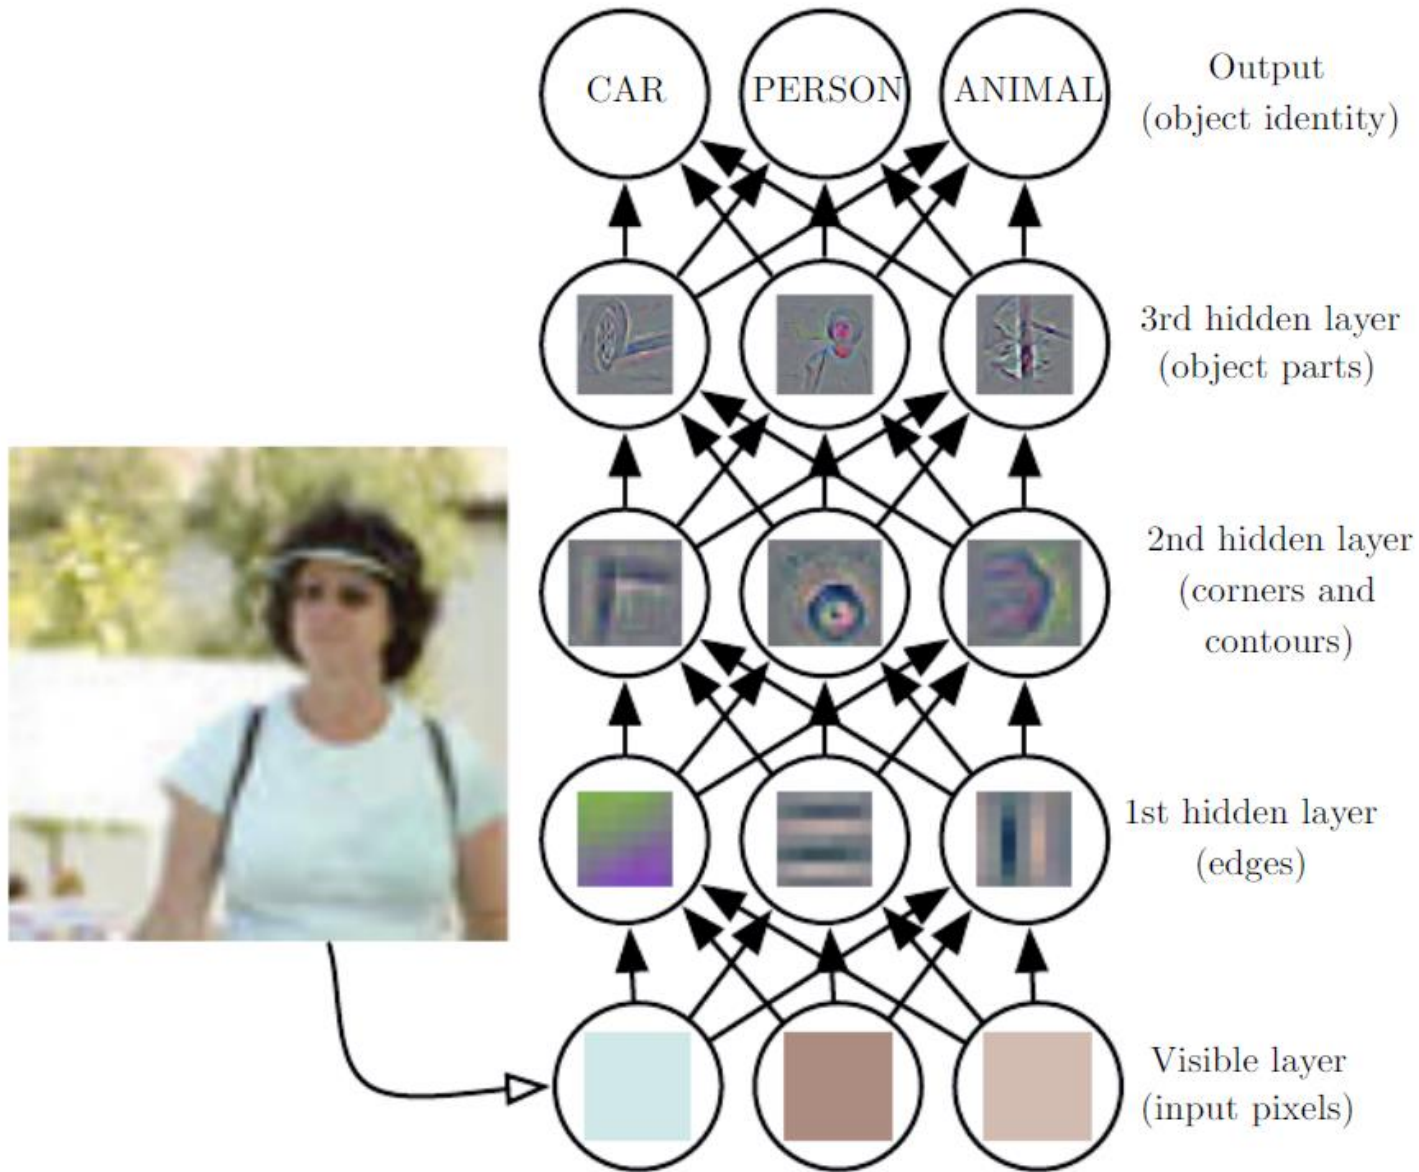
- How do we recognize such “eye”?
- We can divide the image into smaller pieces
- We will try to **match** (convolute) each piece with a pattern, called filter using convolution



F23	F78	F45	F22	F1
F19	F15	F16	F14	F5

If we see this pattern -> This is an eye!

Intuition for Pattern Recognition Using NN





- If we want to detect if this is a person:
 - Detect the edges
 - From those edges, we form corners and complicated contours
 - Those corners and contours are combined to form objects, e.g., eyes, noses
 - Those objects then form a bigger objects, e.g, heads
 - Those bigger objects then form more bigger object, e.g., body
 - ...
- Higher layers will combine smaller features from lower layers to recognize more complicated objects

Convolution Layer

- 32x32x3 image is convolved with a 5x5x3 filter


Convolution: Slide the filter over the image, at each point we compute the dot product of the filter with the region covered by the image



w : weight of the filter
 x : sub-image covered by the filter
 b : bias


Convolution Layer

- Convolve (slide) the filter over the entire image, the resulting dot products at each location create a new matrix called activation map



Convolution Layer


- Each filter creates one activation map



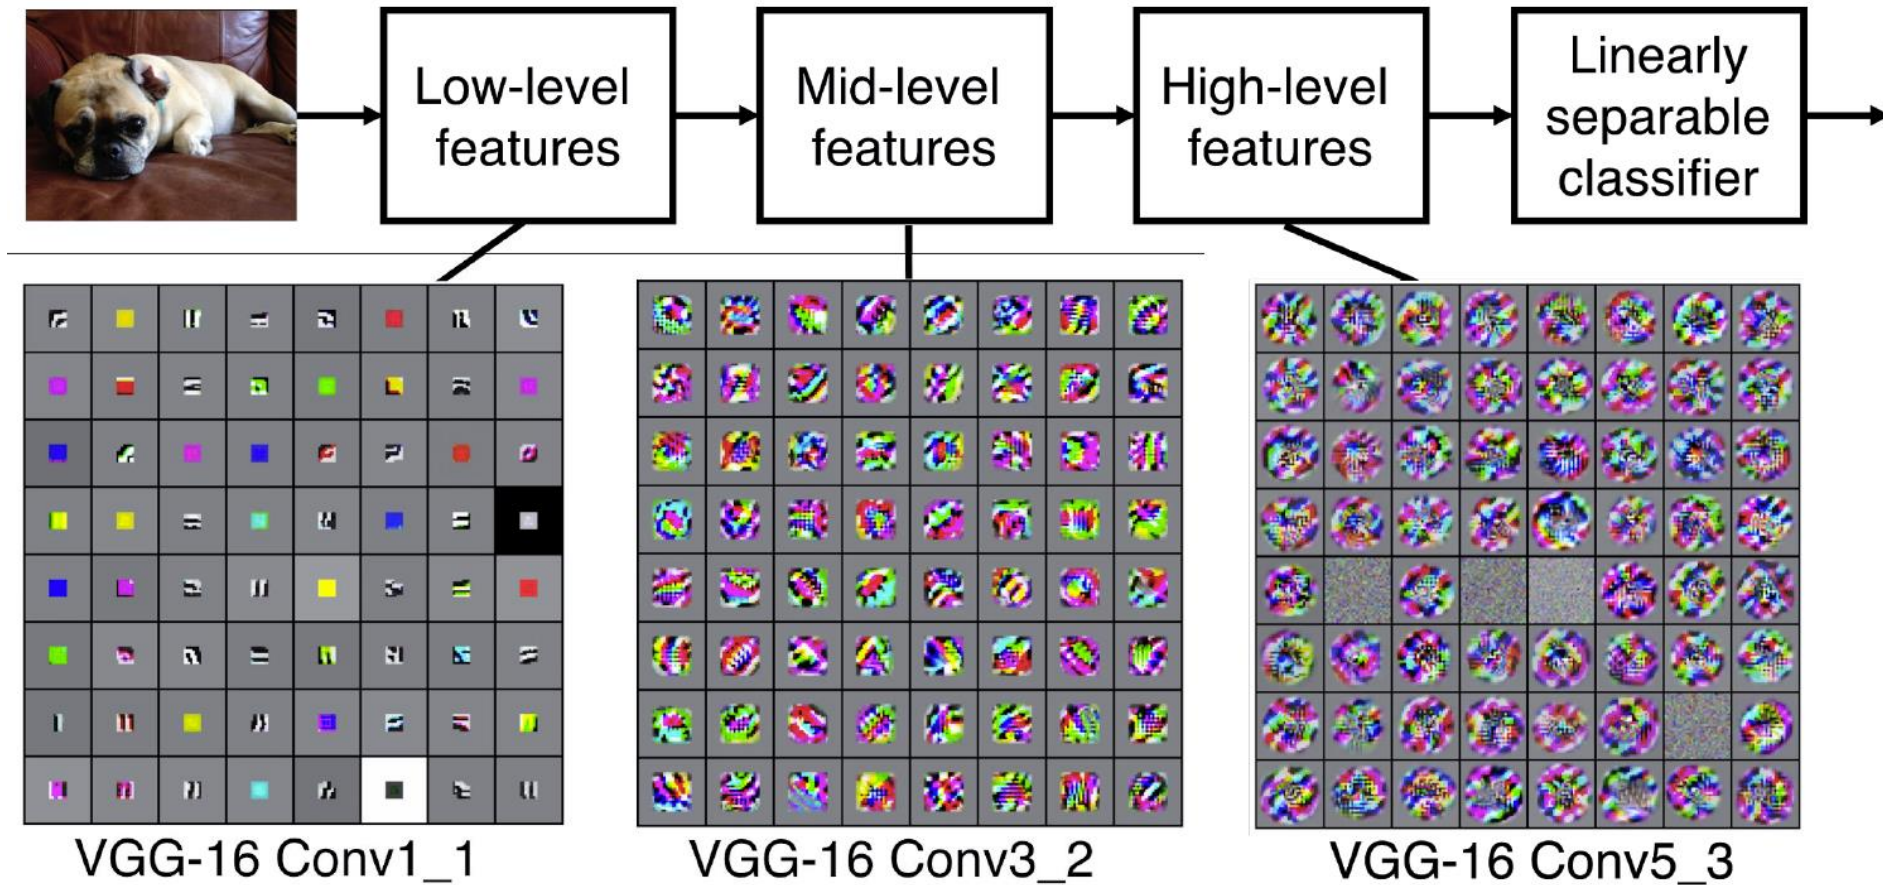
Newly formed activation maps are stacked up to create new image

Example: ConvNet Architecture

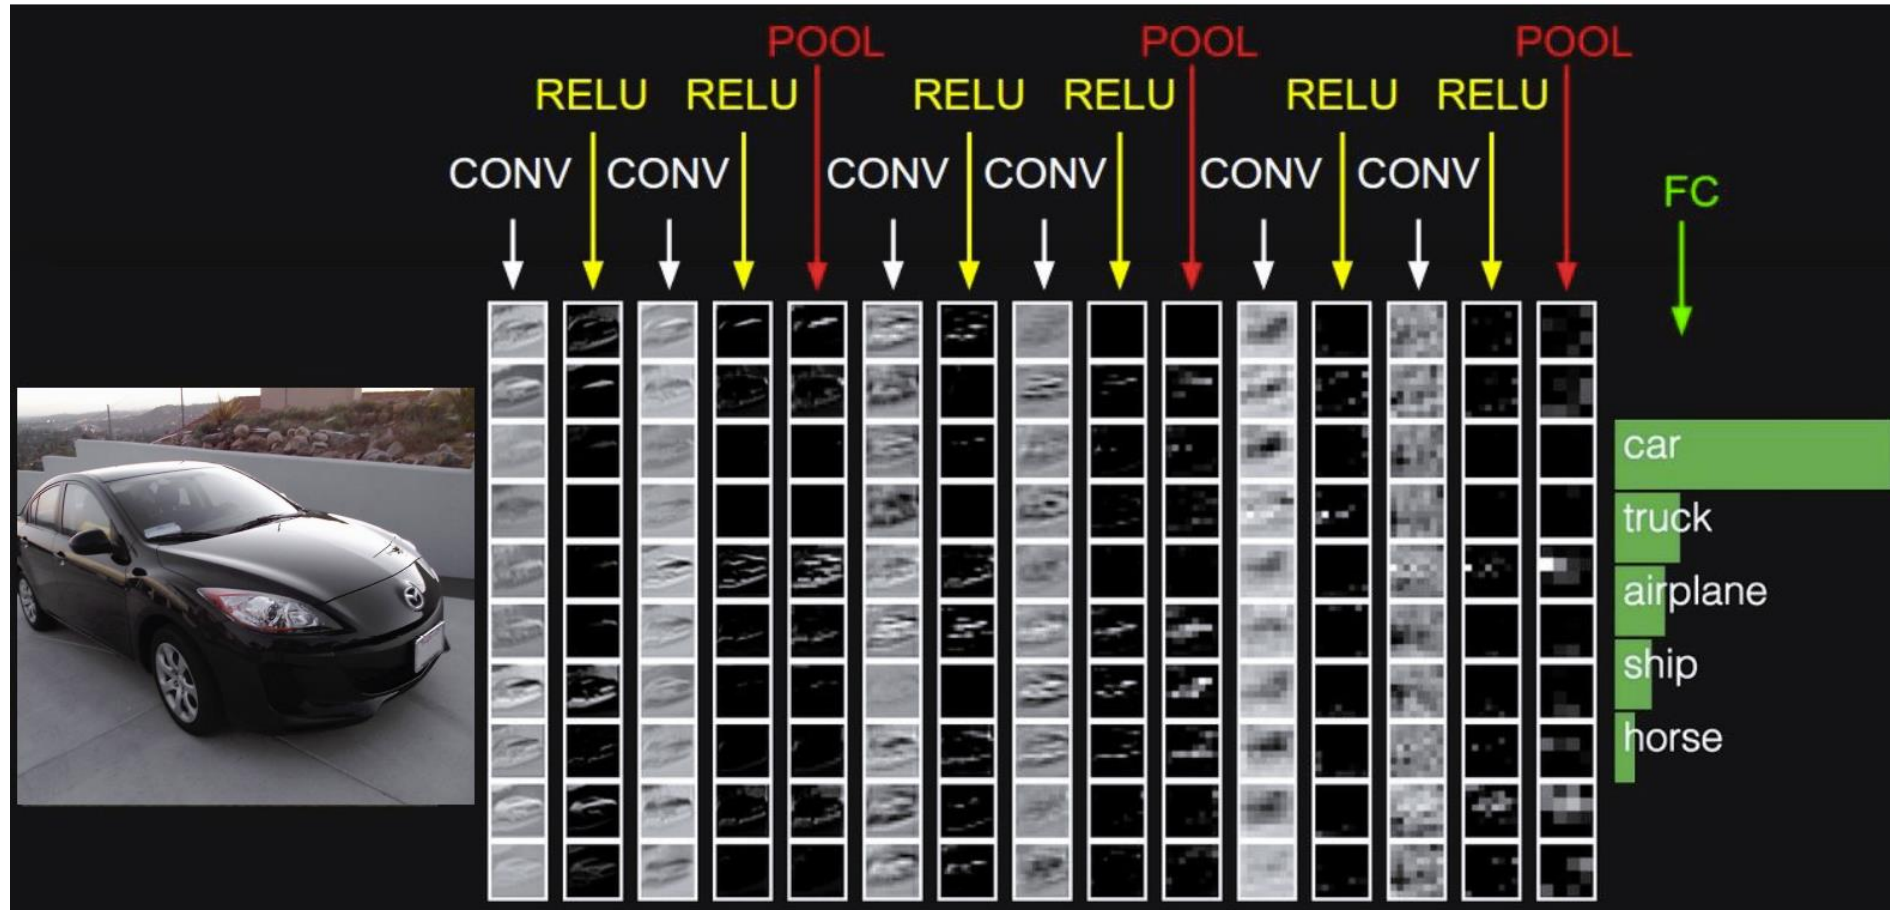
- A sequence of convolution layers




VGG-16 Architecture



Car Recognition Example




Convolution in Details




Convolution with **stride 1**

Convolution in Details



Convolution with **stride 2**

Padding



Insert 0 outside of the border of the image
7x7 image padded with one more pixel at the border
then convolved with one 3x3 filter will produce 7x7 output

Convolution Example

Input Volume (+pad 1) (7x7x3)

$x[:, :, 0]$

0	0	0	0	0	0	0
0	2	2	1	0	2	0
0	1	1	0	0	0	0
0	0	2	0	1	0	0
0	1	0	2	2	0	0
0	1	0	1	2	2	0
0	0	0	0	0	0	0

$x[:, :, 1]$

0	0	0	0	0	0	0
0	0	0	2	1	1	0
0	1	0	2	0	2	0
0	1	0	1	2	1	0
0	1	2	1	0	0	0
0	0	0	1	2	0	0
0	0	0	0	0	0	0

$x[:, :, 2]$

0	0	0	0	0	0	0
0	0	0	0	2	0	0
0	1	0	2	2	0	0
0	0	2	2	1	1	0
0	0	2	2	0	2	0
0	0	2	0	0	2	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

$w0[:, :, 0]$

1	0	0
0	0	0
0	1	-1

$w0[:, :, 1]$

0	0	0
-1	-1	1
0	-1	1

$w0[:, :, 2]$

0	1	0
-1	-1	0
-1	-1	1

Bias b0 (1x1x1)

$b0[:, :, 0]$

1

Filter W1 (3x3x3)

$w1[:, :, 0]$

0	-1	-1
-1	0	-1
-1	0	-1

$w1[:, :, 1]$

-1	-1	0
1	1	1
-1	1	1

$w1[:, :, 2]$

-1	1	-1
1	-1	0
1	0	-1

Bias b1 (1x1x1)

$b1[:, :, 0]$

0

Output Volume (3x3x2)

$o[:, :, 0]$

-1	-2	-7
5	-4	-6
1	2	1

$o[:, :, 1]$


-2	0	8
-2	-3	-4
-4	-4	0

toggle movement

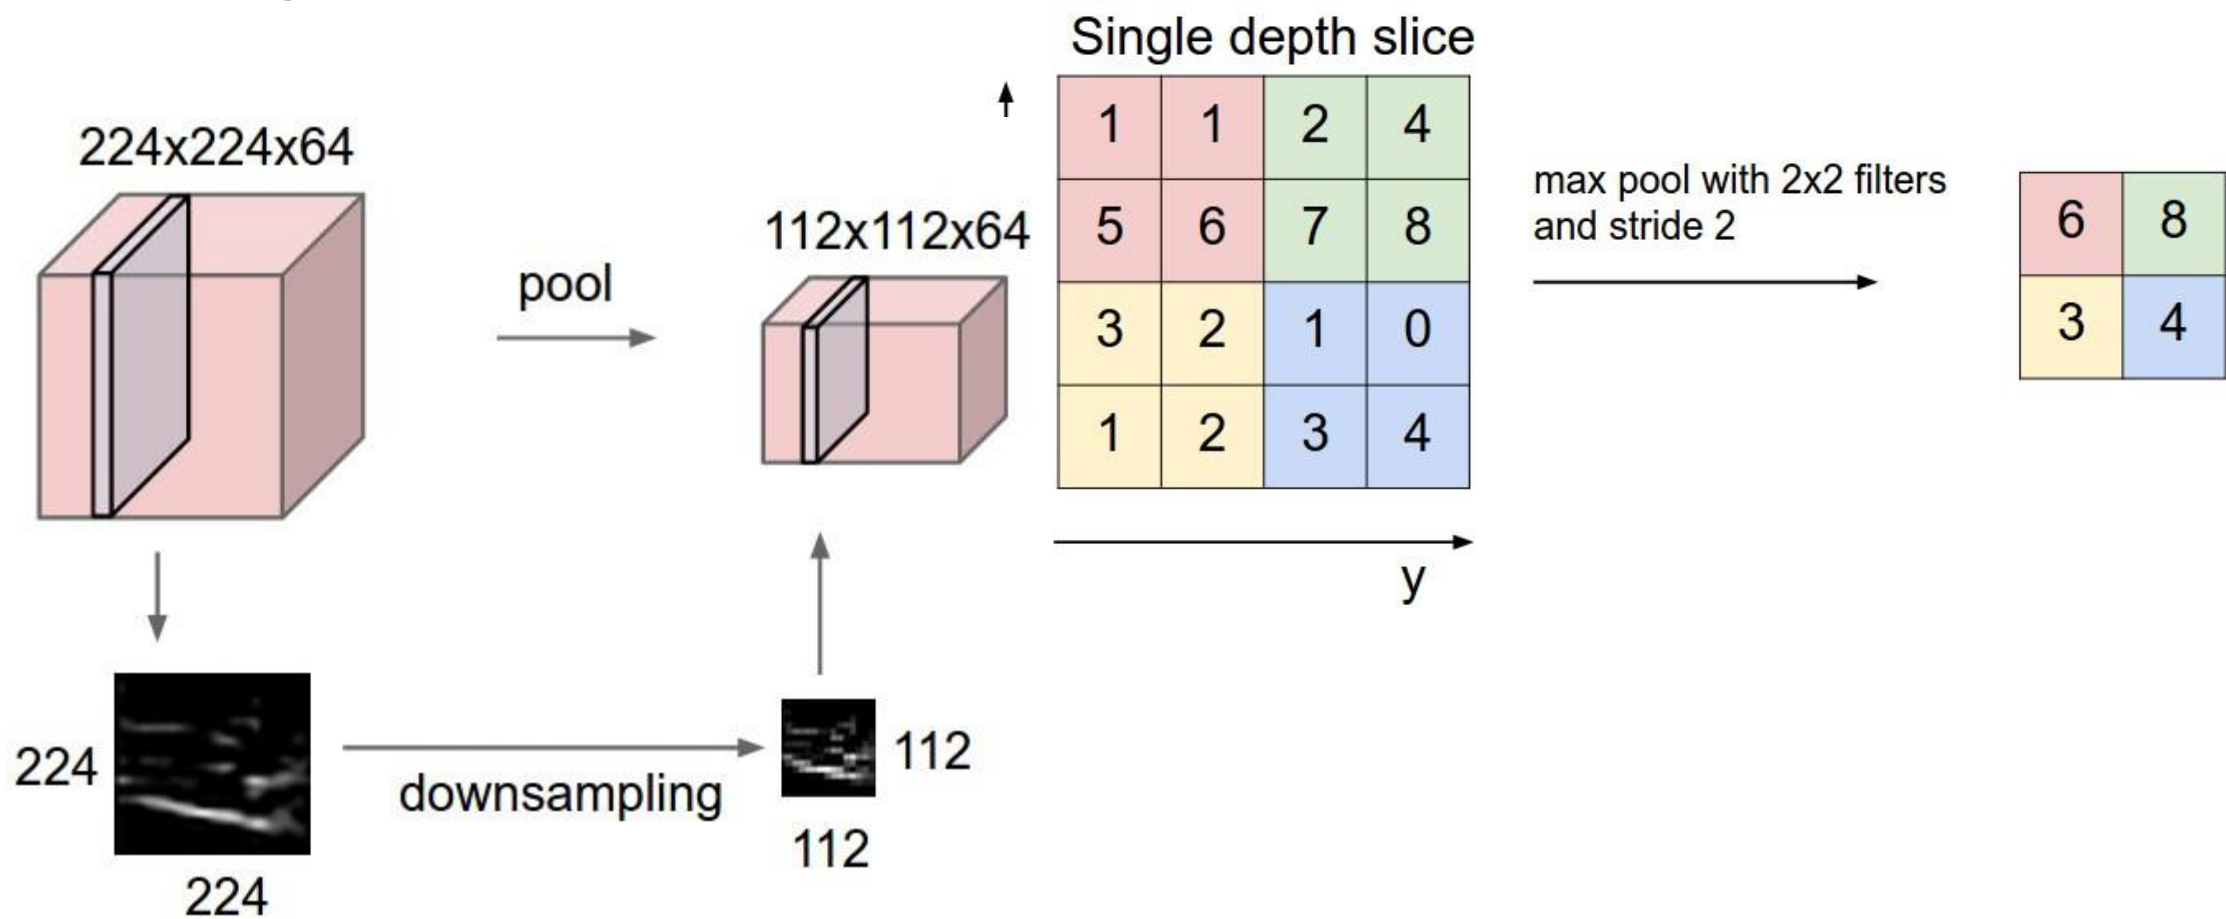
toggle movement

Pooling

- What happens if we have too many layers and the activation maps are big?
 - We want to scale it down
- How do we scale it down?
 - For each submatrix, say 3x3 as in this case, we replace it with a single value
- Max pooling: submatrix is replaced with its maximum value, in this case is 9
- Average pooling: submatrix is replaced with its average value, in this case is 7




Pooling



Convolution Questions

- What happens if we expand w so that it has the same size as the image? (effectively, we use a fully connected network for convolution)
 - It is very expensive to do the computation when the image is large
- Should we expand w so that it has the same size as the image?
 - Perhaps it is not necessary because image pixels are more related to their nearby pixels

Convolution: Slide the filter over the image, at each point we compute the dot product of the filter with the region covered by the image



CNN Practice